

azure-archive-solution

A ready-to-upload GitHub repo for the **Cost Optimization — Azure serverless + Cosmos DB** assignment. Drop the files below into a repository named `azure-archive-solution` and push to GitHub.

Repo structure

```
azure-archive-solution/
├─ README.md
├─ diagrams/
│   └─ architecture.mmd
├─ src/
│   ├── archive_function/
│   │   ├── function_app.py
│   │   └─ requirements.txt
│   └─ api/
│       └─ read_handler.py
├─ infra/
│   └─ azurecli-commands.sh
├─ docs/
│   ├── failure_modes.md
│   └─ cost_strategy.md
├─ sample_data/
│   └─ sample-record.json
└─ LICENSE
```

Files (copy each file content into that path)

README.md

```
# Azure Archive Solution

This repository contains a working solution for the Cost Optimization: Managing Billing Records in Azure Serverless Architecture assignment.

## Overview

- Move cold data (>90 days) from Cosmos DB into compressed blobs in Azure Blob Storage.
```

- Keep a small stub in Cosmos DB so existing API contracts remain unchanged.
- Use an Azure Function to copy-verify-replace documents.
- API reads Cosmos; if stub found, fetches the full record from Blob Storage and caches in Redis.

What to upload

This repo contains the implementation for:

- Archival Azure Function (Python)
- API read handler (Python)
- Infra scripts (Azure CLI)
- Architecture diagram (Mermaid)
- Docs: failure modes and cost strategy

Deploy (high-level)

1. Create resources (see ``infra/azurecli-commands.sh``).
2. Configure environment variables / Managed Identity for the Function App.
3. Deploy the Function App and API.
4. Run the archival job manually once to bootstrap, then enable Timer trigger.

Environment variables (example)

Set these in Function App configuration or use Managed Identity + Key Vault:

- ``COSMOS_ENDPOINT``
- ``COSMOS_KEY`` or use AAD/Managed Identity
- ``COSMOS_DB``
- ``COSMOS_CONTAINER``
- ``BLOB_CONNECTION_STRING`` or use Managed Identity
- ``BLOB_CONTAINER``
- ``REDIS_CONNECTION_STRING`` (optional)

Files of interest

- ``src/archive_function/function_app.py`` – the archival worker
- ``src/api/read_handler.py`` – read-time handler with blob fallback
- ``infra/azurecli-commands.sh`` – quick infra bootstrap script

Notes

- ****Do not**** commit secrets to the repo. Use GitHub Secrets / Azure Key Vault.
- This repo uses gzip JSON for archive objects. You can switch to Parquet for analytics.

```

## diagrams/architecture.mmd

```mermaid
flowchart LR
 Client[Client / API] --> API[API]
 API --> Cosmos[Cosmos DB]
 Cosmos -->|archival job reads old docs| ArchFn[Azure Function (Archive)]
 ArchFn --> Blob[Azure Blob Storage]
 ArchFn --> Cosmos[Update doc -> stub]
 API -->|if stub| BlobFetch[Blob fetch + Cache]
 BlobFetch --> Blob
 BlobFetch --> Redis[Redis Cache]

```

## src/archive\_function/function\_app.py

```

import os
import json
import gzip
import hashlib
import logging
from datetime import datetime, timedelta, timezone

import azure.functions as func
from azure.storage.blob import BlobServiceClient
from azure.cosmos import CosmosClient, PartitionKey

Environment configuration (use Function App settings or Key Vault)
COSMOS_ENDPOINT = os.environ.get("COSMOS_ENDPOINT")
COSMOS_KEY = os.environ.get("COSMOS_KEY")
COSMOS_DB = os.environ.get("COSMOS_DB", "billing")
COSMOS_CONTAINER = os.environ.get("COSMOS_CONTAINER", "records")
BLOB_CONN_STR = os.environ.get("BLOB_CONNECTION_STRING")
BLOB_CONTAINER = os.environ.get("BLOB_CONTAINER", "billing-archive")

Archival threshold in days
ARCHIVE_DAYS = int(os.environ.get("ARCHIVE_DAYS", "90"))
BATCH_SIZE = int(os.environ.get("ARCHIVE_BATCH_SIZE", "100"))

Initialize clients
cosmos_client = CosmosClient(COSMOS_ENDPOINT, COSMOS_KEY)
db_client = cosmos_client.get_database_client(COSMOS_DB)
container = db_client.get_container_client(COSMOS_CONTAINER)

blob_service = BlobServiceClient.from_connection_string(BLOB_CONN_STR)
archive_container = blob_service.get_container_client(BLOB_CONTAINER)

```

```

def md5_bytes(b: bytes) -> str:
 return hashlib.md5(b).hexdigest()

def make_blob_path(doc: dict) -> str:
 created = doc.get("createdAt", datetime.now(timezone.utc).isoformat())
 # Expect ISO format, fallback safe-parsing simple slicing
 year = created[:4]
 month = created[5:7] if len(created) >= 7 else "01"
 return f"year={year}/month={month}/{doc['id']}.json.gz"

def upload_blob_if_missing(path: str, data: bytes):
 blob = archive_container.get_blob_client(path)
 if not blob.exists():
 blob.upload_blob(data, overwrite=False)
 return blob

def archive_document(doc: dict):
 # 1) serialize and compress
 payload = json.dumps(doc, default=str).encode("utf-8")
 gz = gzip.compress(payload)
 blob_path = make_blob_path(doc)

 # 2) upload
 blob = upload_blob_if_missing(blob_path, gz)

 # 3) verify checksum by re-downloading (simple verification step)
 downloaded = blob.download_blob().readall()
 if md5_bytes(downloaded) != md5_bytes(gz):
 raise RuntimeError("Checksum mismatch after upload")

 # 4) replace doc with stub (preserve id and partition key)
 stub = {
 "id": doc["id"],
 "partitionKey": doc.get("partitionKey", doc.get("pk", "default")),
 "archived": True,
 "archiveLocation": blob_path,
 "archiveSize": len(gz),
 "archivedAt": datetime.now(timezone.utc).isoformat(),
 # keep a subset of searchable fields to avoid rehydrating for simple
 queries
 "searchable": {
 "invoiceNo": doc.get("invoiceNo"),
 "amount": doc.get("amount"),

```

```

 "createdAt": doc.get("createdAt")
 }
}

Upsert stub (this will overwrite the existing doc)
container.upsert_item(stub)

def main(mytimer: func.TimerRequest) -> None:
 logging.info("Archive function started")
 cutoff = (datetime.now(timezone.utc) -
timedelta(days=ARCHIVE_DAYS)).isoformat()

 query = (
 "SELECT * FROM c WHERE c.createdAt < @cutoff AND (NOT
IS_DEFINED(c.archived) OR c.archived = false)"
)
 parameters = [{"name": "@cutoff", "value": cutoff}]

 # Query with pagination and batching
 items_iterable = container.query_items(
 query=query,
 parameters=parameters,
 enable_cross_partition_query=True,
)

 count = 0
 for item in items_iterable:
 try:
 archive_document(item)
 count += 1
 if count >= BATCH_SIZE:
 logging.info(f"Processed batch of {count}, stopping to avoid RU
spikes")
 break
 except Exception as e:
 logging.exception(f"Failed to archive {item.get('id')}: {e}")
 # leave the item for retry

 logging.info(f"Archive run complete. Archived {count} documents.")

```

## src/archive\_function/requirements.txt

```
azure-functions
azure-storage-blob
azure-cosmos
```

## src/api/read\_handler.py

```
import os
import json
import gzip

from azure.cosmos import CosmosClient
from azure.storage.blob import BlobServiceClient

Optional: Redis caching (not implemented here) – placeholder for integration

COSMOS_ENDPOINT = os.environ.get("COSMOS_ENDPOINT")
COSMOS_KEY = os.environ.get("COSMOS_KEY")
COSMOS_DB = os.environ.get("COSMOS_DB", "billing")
COSMOS_CONTAINER = os.environ.get("COSMOS_CONTAINER", "records")
BLOB_CONN_STR = os.environ.get("BLOB_CONNECTION_STRING")
BLOB_CONTAINER = os.environ.get("BLOB_CONTAINER", "billing-archive")

cosmos_client = CosmosClient(COSMOS_ENDPOINT, COSMOS_KEY)
container =
cosmos_client.get_database_client(COSMOS_DB).get_container_client(COSMOS_CONTAINER)
blob_service = BlobServiceClient.from_connection_string(BLOB_CONN_STR)
archive_container = blob_service.get_container_client(BLOB_CONTAINER)

def get_record(id: str, partition_key: str) -> dict:
 doc = container.read_item(item=id, partition_key=partition_key)
 if doc.get("archived"):
 # Fetch blob
 blob_path = doc["archiveLocation"]
 blob = archive_container.get_blob_client(blob_path)
 data = blob.download_blob().readall()
 payload = gzip.decompress(data)
 record = json.loads(payload)
 return record
 return doc

Example usage (for local testing)
```

```

if __name__ == "__main__":
 import sys
 id = sys.argv[1]
 pk = sys.argv[2]
 print(get_record(id, pk))

```

## infra/azurecli-commands.sh

```

#!/bin/bash
set -e

RG="rg-billing-$(date +%s)"
LOCATION="eastus"
STORAGE_ACCOUNT="sbillingarchive$RANDOM"
COSMOS_ACCOUNT="cosmosbilling$RANDOM"
FUNCTION_APP="fn-archive-$RANDOM"

1) Resource group
az group create -n $RG -l $LOCATION

2) Storage account for Function & Blobs
az storage account create -n $STORAGE_ACCOUNT -g $RG -l $LOCATION --sku
Standard_LRS

Create blob container
az storage container create --name billing-archive --account-name
$STORAGE_ACCOUNT

3) Cosmos DB account (serverless)
az cosmosdb create -n $COSMOS_ACCOUNT -g $RG --capabilities EnableServerless

4) Function App (Consumption plan) - requires a storage account
az functionapp create --resource-group $RG --consumption-plan-location
$LOCATION \
 --name $FUNCTION_APP --storage-account $STORAGE_ACCOUNT --runtime python --
runtime-version 3.11

Print outputs – remember to set these as config for your function
az cosmosdb keys list -n $COSMOS_ACCOUNT -g $RG --type keys

echo "Resource group: $RG"
echo "Storage account: $STORAGE_ACCOUNT"
echo "Cosmos account: $COSMOS_ACCOUNT"
echo "Function App: $FUNCTION_APP"

```

## docs/failure\_modes.md

```
Failure Modes & Mitigations

1. Partial upload / corrupted blob
 - Verify checksum after upload. If mismatch, abort and retry.

2. Function crashes mid-run
 - Make steps idempotent. Use blob existence and upsert semantics to resume.

3. Read while migrating
 - Use copy-then-replace. Reads will see original doc or a stub that points to a valid blob.

4. RU throttling on Cosmos
 - Batch archival, add sleeps, and use continuation tokens. Monitor RU and back off.

5. Large documents rehydration latency
 - Add Redis cache and/or return lightweight metadata while streaming the payload.
```

## docs/cost\_strategy.md

```
Cost Optimization Strategy

- Move cold objects (>90 days) to Blob Storage (Cool / Archive tiers) to save on storage costs.
- Keep only minimal searchable metadata in Cosmos to reduce RU consumption.
- Use serverless or autoscale Cosmos to avoid fixed RU costs.
- Batch archive operations to reduce transaction overhead.
- Consider Parquet for analytics and smaller read IO when doing large-scale analytics.
```

## sample\_data/sample-record.json

```
{
 "id": "sample-123",
 "partiti
```